

Business Understanding

Problem Statement

Developing a classifier to predict whether a customer will soon stop doing business with SyriaTel, a telecommunications company. This binary classification task aims to identify patterns in customer behavior and demographic information that may indicate a likelihood to churn. The ultimate goal is to assist SyriaTel in reducing the financial impact of customer churn by implementing proactive retention strategies. We further aim to create a reliable churn rate prediction model by thoroughly examining important features based on the historical company's data

Column Name Description

- **state:** the state the user lives in
- **account length:** the number of days the user has this account
- **area code:** the code of the area the user lives in
- **phone number:** the phone number of the user
- **international plan:** true if the user has the international plan, otherwise false
- **voice mail plan:** true if the user has the voice mail plan, otherwise false
- **number vmail messages:** the number of voice mail messages the user has sent
- **total day minutes:** total number of minutes the user has been in calls during the day
- **total day calls:** total number of calls the user has done during the day
- **total day charge:** total amount of money the user was charged by the Telecom company for calls during the day
- **total eve minutes:** total number of minutes the user has been in calls during the evening
- **total eve calls:** total number of calls the user has done during the evening
- **total eve charge:** total amount of money the user was charged by the Telecom company for calls during the evening
- **total night minutes:** total number of minutes the user has been in calls during the night
- **total night calls:** total number of calls the user has done during the night
- **total night charge:** total amount of money the user was charged by the Telecom company for calls during the night
- **total intl minutes:** total number of minutes the user has been in international calls
- **total intl calls:** total number of international calls the user has done
- **total intl charge:** total amount of money the user was charged by the Telecom company for international calls
- **customer service calls:** number of customer service calls the user has done
- **churn:** true if the user terminated the contract, otherwise false

Research questions

1. what are the indicators that show likely to churn?

2. Which state has the highest churn rate

3. Does the type of customer service affect churn?

4. What is the correlation between minutes and churn?

5. Does the presence of an international calls affect the churn rate?

6. Does the account length affect the churn?

7. Does charge affect churn?

Data Understanding

In [1]: `1 pip install xgboost`

Requirement already satisfied: xgboost in c:\users\tabit\anaconda3\lib\site-packages (2.0.3)Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: numpy in c:\users\tabit\anaconda3\lib\site-packages (from xgboost) (1.24.3)

Requirement already satisfied: scipy in c:\users\tabit\anaconda3\lib\site-packages (from xgboost) (1.11.1)

In [2]: `1 # importing libraries
2 import pandas as pd
3 import numpy as np
4 import seaborn as sns
5 import matplotlib.pyplot as plt
6 %matplotlib inline`

In [3]: `1 # Load the data
2 df= pd.read_csv('customer_churn.csv')`

In [4]:

1

checking the first 10 rows

2

df.head(10)

Out[4]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...
5	AL	118	510	391-8027	yes	no	0	223.4	98	37.98	...
6	MA	121	510	355-9993	no	yes	24	218.2	88	37.09	...
7	MO	147	415	329-9001	yes	no	0	157.0	79	26.69	...
8	LA	117	408	335-4719	no	no	0	184.5	97	31.37	...
9	WV	141	415	330-8173	yes	yes	37	258.6	84	43.96	...

10 rows × 21 columns

In [5]:

1

checking the last 10 rows

2

df.tail(10)

Out[5]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge
3323	IN	117	415	362-5899	no	no	0	118.4	126	20.13
3324	WV	159	415	377-1164	no	no	0	169.8	114	28.87
3325	OH	78	408	368-8555	no	no	0	193.4	99	32.88
3326	OH	96	415	347-6812	no	no	0	106.6	128	18.12
3327	SC	79	415	348-3830	no	no	0	134.7	98	22.90
3328	AZ	192	415	414-4276	no	yes	36	156.2	77	26.55
3329	WV	68	415	370-3271	no	no	0	231.1	57	39.29
3330	RI	28	510	328-8230	no	no	0	180.8	109	30.74
3331	CT	184	510	364-6381	yes	no	0	213.8	105	36.35
3332	TN	74	415	400-4344	no	yes	25	234.4	113	39.85

10 rows × 21 columns



In [6]:

1

checking the shape of the data

2

df.shape

Out[6]: (3333, 21)

In [7]:

1

description of the data

2

df.describe()

Out[7]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total day charge
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.900000
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.700000
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000



In [8]:

```
1 # Information about the data
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                        3333 non-null   object
4   international plan                  3333 non-null   object
5   voice mail plan                     3333 non-null   object
6   number vmail messages               3333 non-null   int64
7   total day minutes                   3333 non-null   float64
8   total day calls                     3333 non-null   int64
9   total day charge                     3333 non-null   float64
10  total eve minutes                    3333 non-null   float64
11  total eve calls                      3333 non-null   int64
12  total eve charge                     3333 non-null   float64
13  total night minutes                  3333 non-null   float64
14  total night calls                    3333 non-null   int64
15  total night charge                   3333 non-null   float64
16  total intl minutes                   3333 non-null   float64
17  total intl calls                     3333 non-null   int64
18  total intl charge                    3333 non-null   float64
19  customer service calls               3333 non-null   int64
20  churn                                3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

Data Preparation

In [9]:

```
1 # Convert column names to lowercase and replace empty spaces with under
2 df.columns = df.columns.str.lower().str.replace(' ', '_')
```

In [10]:

```
1 # Replace false/true with 0,1
2 df['churn'] = df['churn'].replace({False:0, True:1})
3 df.sample(3)
```

Out[10]:

	state	account_length	area_code	phone_number	international_plan	voice_mail_plan	n
2214	CT	90	415	347-6994	no	no	
718	AK	127	408	383-9255	no	no	
3162	UT	81	415	355-6422	no	no	

3 rows × 21 columns



```
In [11]: 1 # checking for null values
         2 df.isnull().sum()
```

```
Out[11]: state                                0
         account_length                       0
         area_code                           0
         phone_number                         0
         international_plan                   0
         voice_mail_plan                     0
         number_vmail_messages               0
         total_day_minutes                   0
         total_day_calls                     0
         total_day_charge                    0
         total_eve_minutes                   0
         total_eve_calls                     0
         total_eve_charge                    0
         total_night_minutes                 0
         total_night_calls                   0
         total_night_charge                  0
         total_intl_minutes                  0
         total_intl_calls                    0
         total_intl_charge                   0
         customer_service_calls              0
         churn                               0
         dtype: int64
```

```
In [12]: 1 #checking for duplicates
         2 df.duplicated().sum()
```

```
Out[12]: 0
```

EDA

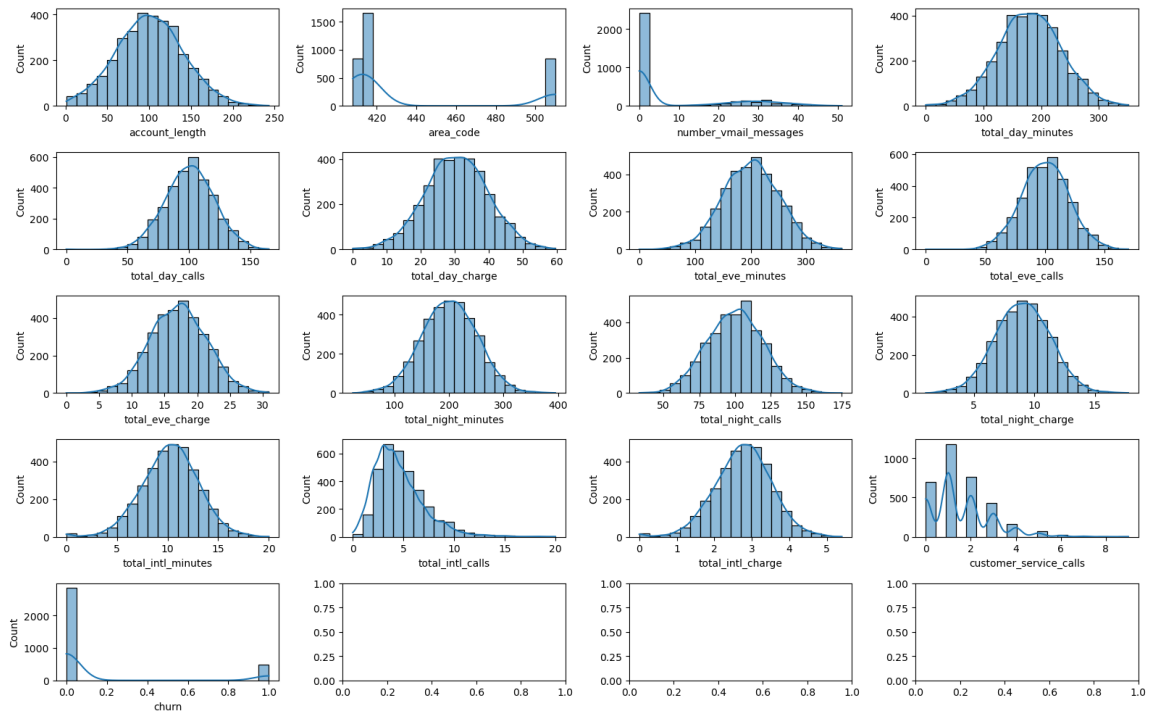
Distribution of Numeric Features in the Dataset

```
In [13]: 1 # checking for data types
         2 df2 = df.select_dtypes(include=['float64', 'int64'])
```

```

In [14]: 1 fig, axes = plt.subplots(nrows=5, ncols=4, figsize=(16, 10))
2
3 # Flatten the axes for easy iteration
4 axes = axes.flatten()
5
6 # Iterate through numeric columns and plot histograms
7 for i, feature in enumerate(df2):
8     sns.histplot(df[feature], ax=axes[i], kde=True, bins=20)
9     axes[i].set_xlabel(feature)
10    axes[i].set_ylabel('Count')
11
12 plt.tight_layout()
13 plt.show()


```




Subscription Plan Distribution

```
In [15]: 1 plt.figure(figsize=(8, 6))
2
3 # Plot the first pie chart
4 plt.subplot(1, 2, 1)
5 ax1 = df['international_plan'].value_counts()
6 plt.pie(ax1, labels=['No', 'Yes'], autopct='%0f%%', startangle=90)
7 plt.title('International Plan Subscription Distribution')
8
9 # Plot the second pie chart
10 plt.subplot(1, 2, 2)
11 ax2 = df['voice_mail_plan'].value_counts()
12 plt.pie(ax2, labels=['No', 'Yes'], autopct='%0f%%', startangle=90)
13 plt.title('Voice Mail Plan Distribution')
14
15 # Adjust layout for better spacing
16 plt.tight_layout()
17
18 # Show the plot
19 plt.show()
```

International Plan Subscription Distribution



Voice Mail Plan Distribution




Minutes Distribution

```
In [16]: 1 #function to create a pie chart
2 def pie (title, x_axis, labels):
3     gsdfgdf
4     plt.figure(figsize=(6, 6))
```




```
In [17]: 1 col_sum = ['total_day_minutes', 'total_eve_minutes', 'total_night_minut
2
3 # Calculate the sum for each column
4 sums = df[col_sum].sum()
5
6 # Plot a pie chart
7 plt.figure(figsize=(6, 6))
8 plt.pie(sums, labels=sums.index, autopct='%1.1f%%', startangle=90)
9 plt.title('Total Minutes Distribution')
10 plt.show()
```



- This shows the distribution of total minutes across different call categories.
- Each slice of the pie represents the proportion of total minutes for a specific call category with total evening minutes being the highest.

Churn Distribution

```
In [18]: 1 churn_counts = df['churn'].value_counts()
2 plt.figure(figsize=(6, 6))
3 plt.pie(churn_counts, labels=['No', 'Yes'], autopct='%1.1f%%', startang
4 plt.title('Churn Distribution')
5 plt.show()
```




- The pie chart demonstrates that a smaller percentage of customers are churning (leaving) compared to those who are staying. This insight is valuable for businesses to understand the current state of customer retention and can inform strategies aimed at reducing churn and enhancing customer satisfaction.

Area Code Distribution

```
In [19]: 1 area_code_counts = df['area_code'].value_counts()
2
3 # Plot a pie chart
4 plt.figure(figsize=(6, 6))
5 plt.pie(area_code_counts, labels=area_code_counts.index, autopct='%1.1f%%')
6 plt.title('Distribution of Area Codes')
7 plt.show()
```

Distribution of Area Codes



Churn Distribution by State

In [20]:


```
1 from IPython.display import display
2
3 # Dictionary mapping state initials to full names
4 state_mapping = {
5     'AL': 'Alabama',
6     'AK': 'Alaska',
7     'AZ': 'Arizona',
8     'AR': 'Arkansas',
9     'CA': 'California',
10    'CO': 'Colorado',
11    'CT': 'Connecticut',
12    'DE': 'Delaware',
13    'FL': 'Florida',
14    'GA': 'Georgia',
15    'HI': 'Hawaii',
16    'ID': 'Idaho',
17    'IL': 'Illinois',
18    'IN': 'Indiana',
19    'IA': 'Iowa',
20    'KS': 'Kansas',
21    'KY': 'Kentucky',
22    'LA': 'Louisiana',
23    'ME': 'Maine',
24    'MD': 'Maryland',
25    'MA': 'Massachusetts',
26    'MI': 'Michigan',
27    'MN': 'Minnesota',
28    'MS': 'Mississippi',
29    'MO': 'Missouri',
30    'MT': 'Montana',
31    'NE': 'Nebraska',
32    'NV': 'Nevada',
33    'NH': 'New Hampshire',
34    'NJ': 'New Jersey',
35    'NM': 'New Mexico',
36    'NY': 'New York',
37    'NC': 'North Carolina',
38    'ND': 'North Dakota',
39    'OH': 'Ohio',
40    'OK': 'Oklahoma',
41    'OR': 'Oregon',
42    'PA': 'Pennsylvania',
43    'RI': 'Rhode Island',
44    'SC': 'South Carolina',
45    'SD': 'South Dakota',
46    'TN': 'Tennessee',
47    'TX': 'Texas',
48    'UT': 'Utah',
49    'VT': 'Vermont',
50    'VA': 'Virginia',
51    'WA': 'Washington',
52    'WV': 'West Virginia',
53    'WI': 'Wisconsin',
54    'WY': 'Wyoming'
55 }
56
57
58 grouped_df = df.groupby(['state', 'churn']).size().unstack() # Group by state and churn
59 grouped_df.index = grouped_df.index.map(state_mapping) # Map state initials to full names
60 grouped_df['Total'] = grouped_df.sum(axis=1) # Calculate total number of customers
61
```

```
62 # Display the DataFrame as a table  
63 display(grouped_df)
```

churn	0	1	Total
state			
Alaska	49	3	52
Alabama	72	8	80
Arkansas	44	11	55
Arizona	60	4	64
California	25	9	34
Colorado	57	9	66
Connecticut	62	12	74
NaN	49	5	54
Delaware	52	9	61
Florida	55	8	63
Georgia	46	8	54
Hawaii	50	3	53
Iowa	41	3	44
Idaho	64	9	73
Illinois	53	5	58
Indiana	62	9	71
Kansas	57	13	70
Kentucky	51	8	59
Louisiana	47	4	51
Massachusetts	54	11	65
Maryland	53	17	70
Maine	49	13	62
Michigan	57	16	73
Minnesota	69	15	84
Missouri	56	7	63
Mississippi	51	14	65
Montana	54	14	68
North Carolina	57	11	68
North Dakota	56	6	62
Nebraska	56	5	61
New Hampshire	47	9	56
New Jersey	50	18	68
New Mexico	56	6	62
Nevada	52	14	66
New York	68	15	83
Ohio	68	10	78
Oklahoma	52	9	61
Oregon	67	11	78


churn	0	1	Total
state			
Pennsylvania	37	8	45
Rhode Island	59	6	65
South Carolina	46	14	60
South Dakota	52	8	60
Tennessee	48	5	53
Texas	54	18	72
Utah	62	10	72
Virginia	72	5	77
Vermont	65	8	73
Washington	52	14	66
Wisconsin	71	7	78
West Virginia	96	10	106
Wyoming	68	9	77

```
In [21]: 1 churned_df = df[df['churn'] == 1] # Filter DataFrame to include only churned customers
2 state_churned_count = churned_df['state'].value_counts() # Count the number of churned customers per state
3
4 # Plotting the number of churned customers per state
5 plt.figure(figsize=(15, 6))
6 state_churned_count.plot(kind='bar')
7 plt.title('Number of Churned Customers per State')
8 plt.xlabel('State')
9 plt.ylabel('Number of Churned Customers')
10 plt.xticks(rotation=45)
11 plt.show()
```




- The highest state that did churn was New Jersey state


```
In [22]: 1 churned_df = df[df['churn'] == 0] # Filter DataFrame to include only c
2 state_churned_count = churned_df['state'].value_counts() # Count the n
3
4 # Plotting the number of churned customers per state
5 plt.figure(figsize=(15, 6))
6 state_churned_count.plot(kind='bar')
7 plt.title('Number of Customers who remained per State')
8 plt.xlabel('State')
9 plt.ylabel('Number of Customers who remained')
10 plt.xticks(rotation=45)
11 plt.show()
```



- The highest number of Customers who remained are recorded in West Virginia


```
In [23]: 1 # Plot churn distribution by state
2 df.groupby(['state', 'churn']).size().unstack().plot(kind='bar', figsize=
3
4 plt.title('Churn Distribution by State')
5 plt.xlabel('State')
6 plt.ylabel('Number of Customers')
7 plt.legend(['No', 'Yes'], title='Churn')
8 plt.show()
```



- This shows that West Virginia has the highest number of customers, with a high retention

Correlation of the features

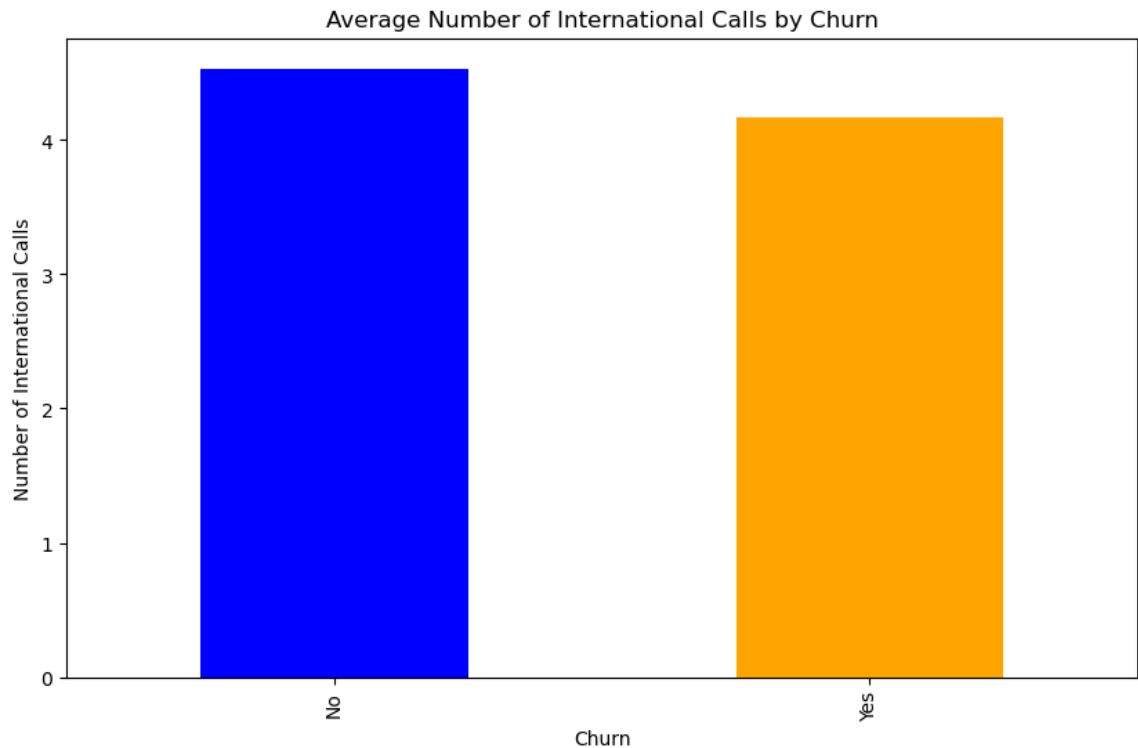
```
In [24]: 1 df2 = df.select_dtypes(include=['float64','int64'])
2 correlation_matrix = df2.corr()
3
4 plt.figure(figsize=(20,7))
5 sns.heatmap(correlation_matrix, annot = True, cmap='coolwarm')
6 plt.show()
```



```
In [25]: 1 corr_matrix = df2.corr()
2 churn_corr = corr_matrix['churn'].sort_values(ascending=False)
3 churn_corr
```


```
Out[25]: churn                1.000000
customer_service_calls        0.208750
total_day_minutes              0.205151
total_day_charge               0.205151
total_eve_minutes              0.092796
total_eve_charge               0.092786
total_intl_charge              0.068259
total_intl_minutes             0.068239
total_night_charge             0.035496
total_night_minutes            0.035493
total_day_calls                0.018459
account_length                 0.016541
total_eve_calls                0.009233
area_code                      0.006174
total_night_calls              0.006141
total_intl_calls               -0.052844
number_vmail_messages          -0.089728
Name: churn, dtype: float64
```

```
In [26]: 1 # Calculate average number of international calls for each churn status
2 total_intl_calls = df.groupby('churn')['total_intl_calls'].mean()
3
4 plt.figure(figsize=(10, 6)) # Create a new figure with specific size
5 total_intl_calls.plot(kind='bar', color=['blue', 'orange']) # Create a
6 plt.title('Average Number of International Calls by Churn') # Add titl
7 plt.xlabel('Churn') # Add Label for x-axis
8 plt.ylabel('Number of International Calls') # Add Label for y-axis
9
10
11 plt.xticks([0, 1], ['No', 'Yes'])
12 plt.show() # Display the plot
13
```



Churn rate by Customer Service Calls

```
In [27]: 1 # Calculate churn rate percentage for each number of customer service c
2 churn_rate = df.groupby('customer_service_calls')['churn'].mean() * 100
3
4 # Plotting a bar plot
5 plt.figure(figsize=(8, 6))
6 ax = sns.barplot(x=churn_rate.index, y=churn_rate.values, palette='flare')
7 ax.bar_label(ax.containers[0], fmt='%.1f%', label_type='edge') # Add %
8
9
10 # Adding title and labels
11 plt.title('Churn Rate Percentage by Customer Service Calls')
12 plt.xlabel('Number of Customer Service Calls')
13 plt.ylabel('Churn Rate %')
14
15 # Display the plot
16 plt.show()
```



There is a positive correlation between the number of customer service calls and the likelihood of churn.

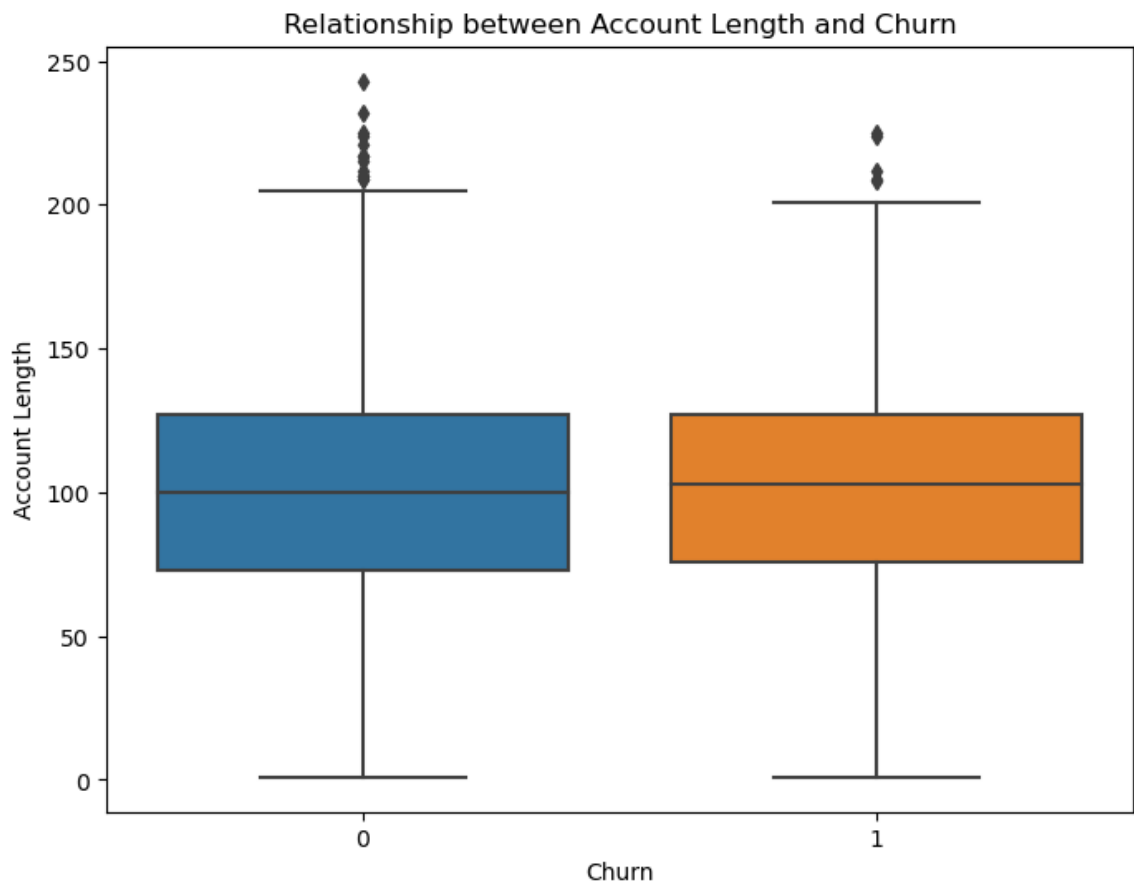
- As the number of customer service calls increases, customers are more likely to churn, or stop using the service. However, there is a disparity in the trend. Specifically, there is a noticeable increase in churn at the 6th customer service call.
- This observation indicates that there might be a particular threshold or point where additional customer service interactions start to have a negative impact on customer retention, potentially leading to a higher churn rate.

- While there is a general trend of increasing churn with more customer service calls, it highlights a potential anomaly or critical point at the 6th customer service call where churn rate spikes, suggesting that this specific interaction might be a key factor in

Relationship between Account Length and Churn

In [28]:

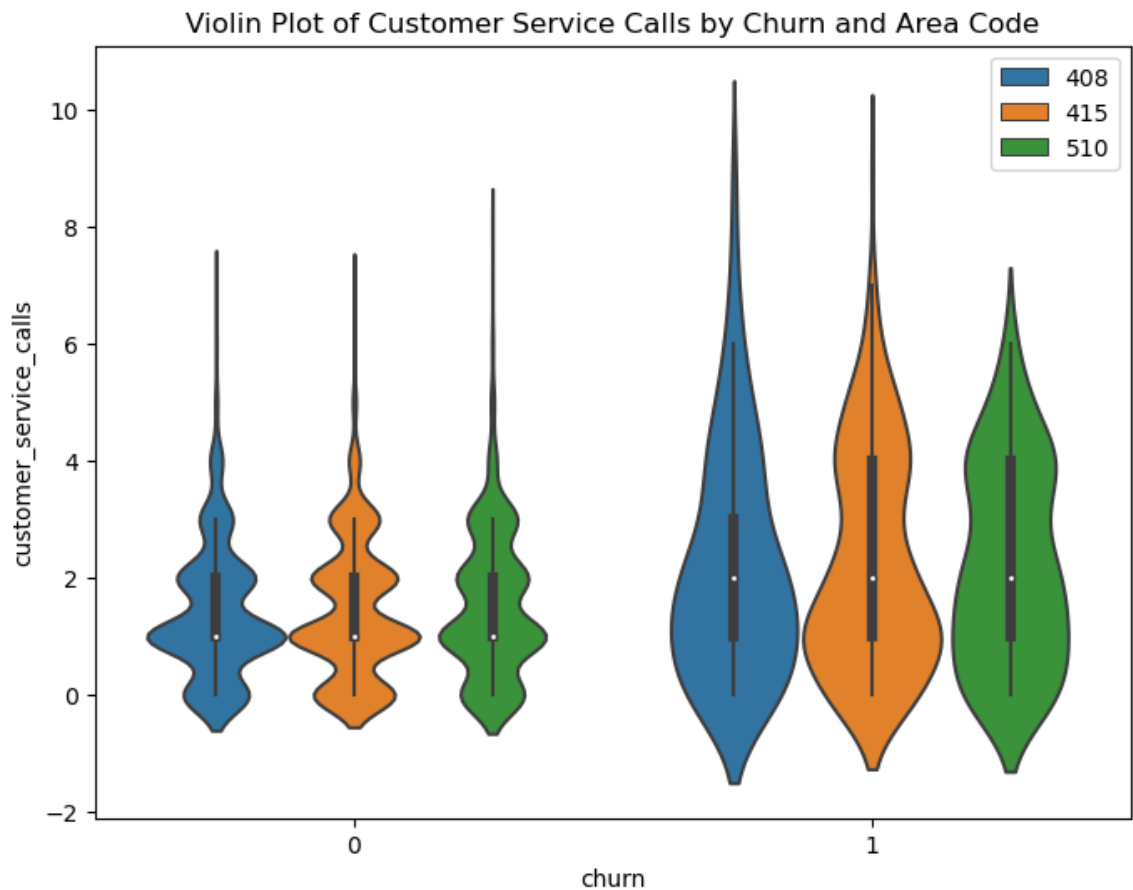
```
1 # Box plot
2 plt.figure(figsize=(8, 6)) # Create a new figure with specific size
3 sns.boxplot(x='churn', y='account_length', data=df) # Create a box plot
4 plt.title('Relationship between Account Length and Churn') # Add title
5 plt.xlabel('Churn') # Add label for x-axis
6 plt.ylabel('Account Length') # Add label for y-axis
7 plt.show() # Display the plot
```



- There's a noticeable difference in median account length between the two groups, it indicates that account length is a factor influencing churn behavior.

Violin plots of the different area codes by customer service calls by churn

```
In [29]: 1 plt.figure(figsize=(8, 6))  
2 sns.violinplot(data=df, x='churn', y='customer_service_calls', hue='are  
3 plt.legend(loc='upper right')  
4 plt.title('Violin Plot of Customer Service Calls by Churn and Area Code  
5 plt.show()
```



- The plot shows the density of data points at different values of customer service calls for each combination of churn status and area code.

In [30]: 1 churn_corr

```
Out[30]: churn                1.000000
customer_service_calls      0.208750
total_day_minutes           0.205151
total_day_charge             0.205151
total_eve_minutes           0.092796
total_eve_charge             0.092786
total_intl_charge            0.068259
total_intl_minutes           0.068239
total_night_charge           0.035496
total_night_minutes          0.035493
total_day_calls              0.018459
account_length               0.016541
total_eve_calls              0.009233
area_code                    0.006174
total_night_calls            0.006141
total_intl_calls             -0.052844
number_vmail_messages        -0.089728
Name: churn, dtype: float64
```

We can drop one of the columns with a 1 correlation (total day minutes and total day charge), (total eve minutes and total eve charge), (total night minutes and total night charge), (total intl minutes and total intl charge) and phone number.

- total day charge
- total eve charge
- total night charge
- total intl charge

```
In [31]: 1 # drop the columns
2 cols_drop = ['total_day_charge', 'total_eve_charge', 'total_night_charg
3 df3 = df.drop(cols_drop, axis=1)
4 df3.head(2)
```

```
Out[31]:
```

	state	account_length	area_code	international_plan	voice_mail_plan	number_vmail_messa
0	KS	128	415	no	yes	
1	OH	107	415	no	yes	

```
In [32]: 1 df3.dtypes
```

```
Out[32]: state                object
account_length             int64
area_code                  int64
international_plan         object
voice_mail_plan            object
number_vmail_messages      int64
total_day_minutes          float64
total_day_calls             int64
total_eve_minutes          float64
total_eve_calls             int64
total_night_minutes        float64
total_night_calls          int64
total_intl_minutes         float64
total_intl_calls           int64
customer_service_calls     int64
churn                      int64
dtype: object
```

```
In [33]: 1 df3['churn'].unique()
```

```
Out[33]: array([0, 1], dtype=int64)
```

```
In [34]: 1 # Dummy variables
2 df4 = pd.get_dummies(df3, columns=['state', 'international_plan', 'voic
```


In [35]:

1	<code>df4.info()</code>
---	-------------------------

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3333 entries, 0 to 3332
```

```
Data columns (total 65 columns):
```

#	Column	Non-Null Count	Dtype
0	account_length	3333 non-null	int64
1	area_code	3333 non-null	int64
2	number_vmail_messages	3333 non-null	int64
3	total_day_minutes	3333 non-null	float64
4	total_day_calls	3333 non-null	int64
5	total_eve_minutes	3333 non-null	float64
6	total_eve_calls	3333 non-null	int64
7	total_night_minutes	3333 non-null	float64
8	total_night_calls	3333 non-null	int64
9	total_intl_minutes	3333 non-null	float64
10	total_intl_calls	3333 non-null	int64
11	customer_service_calls	3333 non-null	int64
12	churn	3333 non-null	int64
13	state_AL	3333 non-null	bool
14	state_AR	3333 non-null	bool
15	state_AZ	3333 non-null	bool
16	state_CA	3333 non-null	bool
17	state_CO	3333 non-null	bool
18	state_CT	3333 non-null	bool
19	state_DC	3333 non-null	bool
20	state_DE	3333 non-null	bool
21	state_FL	3333 non-null	bool
22	state_GA	3333 non-null	bool
23	state_HI	3333 non-null	bool
24	state_IA	3333 non-null	bool
25	state_ID	3333 non-null	bool
26	state_IL	3333 non-null	bool
27	state_IN	3333 non-null	bool
28	state_KS	3333 non-null	bool
29	state_KY	3333 non-null	bool
30	state_LA	3333 non-null	bool
31	state_MA	3333 non-null	bool
32	state_MD	3333 non-null	bool
33	state_ME	3333 non-null	bool
34	state_MI	3333 non-null	bool
35	state_MN	3333 non-null	bool
36	state_MO	3333 non-null	bool
37	state_MS	3333 non-null	bool
38	state_MT	3333 non-null	bool
39	state_NC	3333 non-null	bool
40	state_ND	3333 non-null	bool
41	state_NE	3333 non-null	bool
42	state_NH	3333 non-null	bool
43	state_NJ	3333 non-null	bool
44	state_NM	3333 non-null	bool
45	state_NV	3333 non-null	bool
46	state_NY	3333 non-null	bool
47	state_OH	3333 non-null	bool
48	state_OK	3333 non-null	bool
49	state_OR	3333 non-null	bool
50	state_PA	3333 non-null	bool
51	state_RI	3333 non-null	bool
52	state_SC	3333 non-null	bool
53	state_SD	3333 non-null	bool
54	state_TN	3333 non-null	bool
55	state_TX	3333 non-null	bool

```

56 state_UT          3333 non-null    bool
57 state_VA          3333 non-null    bool
58 state_VT          3333 non-null    bool
59 state_WA          3333 non-null    bool
60 state_WI          3333 non-null    bool
61 state_WV          3333 non-null    bool
62 state_WY          3333 non-null    bool
63 international_plan_yes 3333 non-null    bool
64 voice_mail_plan_yes  3333 non-null    bool
dtypes: bool(52), float64(4), int64(9)
memory usage: 507.9 KB

```

Data Modelling

We start with a Logistic Regression for our baseline model

1. Logistic Regression Model

```

In [36]: 1 # Y Target Variable
          2 y = df4['churn']
          3 X = df4.drop('churn', axis = 1)

```

```

In [37]: 1 # Create Scaler Object to standardize
          2 from sklearn.preprocessing import MinMaxScaler
          3
          4 scaler = MinMaxScaler()
          5
          6 # fit and transform
          7 X_scaled = pd.DataFrame(scaler.fit_transform(X))
          8
          9 X_scaled.head()

```

```

Out[37]:
           0         1         2         3         4         5         6         7         8
0  0.524793  0.068627  0.490196  0.755701  0.666667  0.542755  0.582353  0.595750  0.408451
1  0.438017  0.068627  0.509804  0.460661  0.745455  0.537531  0.605882  0.621840  0.492958
2  0.561983  0.068627  0.000000  0.693843  0.690909  0.333242  0.647059  0.374933  0.500000
3  0.342975  0.000000  0.000000  0.853478  0.430303  0.170195  0.517647  0.467187  0.394366
4  0.305785  0.068627  0.000000  0.475200  0.684848  0.407754  0.717647  0.440290  0.619718

```

5 rows × 64 columns



```
In [38]: 1 # perform train test split
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LogisticRegression
4
5 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
6
7 # Create a logistic regression model
8 logreg = LogisticRegression(fit_intercept=False, C=1e12, solver='liblin
9
10 # fit the model on the training data
11 logreg.fit(X_train, y_train)
```

Out[38]: LogisticRegression(C=1000000000000.0, fit_intercept=False, solver='liblinear')

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [39]: 1 # Generate predictions
2 y_hat_train = logreg.predict(X_train)
3 y_hat_test = logreg.predict(X_test)
```

```
In [40]: 1 # Calculate the performance metrics
2 from sklearn.metrics import accuracy_score, precision_score, recall_score
3
4 # Accuracy
5 accuracy = accuracy_score(y_test, y_hat_test)
6 print(f"Accuracy: {accuracy}")
7
8 # Precision
9 precision = precision_score(y_test, y_hat_test)
10 print(f"Precision: {precision}")
11
12 # Recall
13 recall = recall_score(y_test, y_hat_test)
14 print(f"Recall: {recall}")
15
16 # F1-score
17 f1 = f1_score(y_test, y_hat_test)
18 print(f"F1-score: {f1}")
19
20 # False positive rate(fpr) and true positive rate(tpr)
21 fpr, tpr, thresholds = roc_curve(y_test, y_hat_test)
22
23 # calculate the AUC
24 auc = auc(fpr, tpr)
25 print(f"AUC: {auc}")
```

Accuracy: 0.8515742128935532
Precision: 0.5294117647058824
Recall: 0.1782178217821782
F1-score: 0.26666666666666666
AUC: 0.5749746352727145


- Accuracy: Our classifier shows that our model 85% accurate
- Precision: Out of all the instances our model predicted as positive, approximately 52.94% were actually positive.

- Recall: Being approximately 17.82%, means that out of all the actual positive instances, our model identified approximately 17.82% correctly.
- F1 score provides a balance between precision and recall. In our case, the F1-score is approximately 26.67%.
- AUC is approximately 0.575, suggesting that our model's ability to distinguish between positive and negative classes is more or less the same as random guessing.

In [41]:

```
1 # Build a confusion matrix
2 from sklearn.metrics import confusion_matrix
3
4 # Confusion matrix
5 cm = confusion_matrix(y_test, y_hat_test)
6
7 # Visualize
8 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
9             xticklabels=['Predicted Negative', 'Predicted Positive'],
10            yticklabels=['Actual Negative', 'Actual Positive'])
```

Out[41]: <Axes: >



- The model predicted 18 True Positives, 551 True Negatives, 15 False Positives and 83 False Negatives.
- The model predicted 18 customers would churn and they did.
- The model predicted that 551 customers would not churn and they didn't
- The model predicted that 15 customers would churn but they didn't
- The model incorrectly predicted that 83 customers would not churn but they actually churned

To check if the model is imbalanced

```
In [42]: 1 # Calculate class distribution
2 class_counts = df4['churn'].value_counts()
3 majority_class = class_counts.idxmax()
4 majority_count = class_counts.max()
5
6 # Check for imbalance. Threshold for imbalance --> 0.8
7 if majority_count / len(df4) > 0.8:
8     print('df4 is imbalanced.')
9 else:
10    print('df4 is balanced.')
```

df4 is imbalanced.

- Due to the imbalance of the data we use Random Forest since it handles imbalance better compared to Logistic Regression models

2. Random Forest

```
In [43]: 1 # Create and fit a random forest classifier
2 from sklearn.ensemble import RandomForestClassifier
3
4 clf = RandomForestClassifier(random_state=0)
5 clf = clf.fit(X_train, y_train)
```

```
In [44]: 1 from sklearn.metrics import f1_score, precision_recall_curve, roc_auc_s
2 from sklearn.metrics import roc_curve, auc
3
4 # Predictions on the testing data
5 y_pred = clf.predict(X_test)
6
7 # Calculate accuracy
8 accuracy = accuracy_score(y_test, y_pred)
9 print(f"Accuracy: {accuracy}")
10
11 # Calculate F1-score
12 f1 = f1_score(y_test, y_pred)
13 print(f"F1-score: {f1}")
14
15 # Calculate precision and recall for different thresholds
16 precision, recall, thresholds = precision_recall_curve(y_test, y_pred)
17
18 # Calculate ROC AUC
19 roc_auc = roc_auc_score(y_test, y_pred)
20
21 print(f"ROC AUC: {roc_auc}")
```


Accuracy: 0.9280359820089955

F1-score: 0.6962025316455697

ROC AUC: 0.7705104432704755

- F1-score being 77% is moderately good performance. AUC of 77% is okay for the imbalanced data. This gives a generally good performance of the model.
- Having an accuracy of 92.9%, performs better than that of Logistic Regression Model, 85.3%

```
In [45]: 1 from sklearn.metrics import confusion_matrix
2
3 cm1 = confusion_matrix(y_test, y_pred)
4 # Visualize the confusion matrix
5 sns.heatmap(cm1, annot=True, fmt="d", cmap='Reds',
6             xticklabels=['Predicted Negative', 'Predicted Positive'],
7             yticklabels=['Actual Negative', 'Actual Positive'])
8 plt.show()
```



- The model predicted 55 True Positives, 564 True Negatives, 2 False Positives and 46 False Negatives.
- The model predicted 55 customers would churn and they did.
- The model predicted that 564 customers would not churn and they didn't
- The model predicted that 2 customers would churn but they didn't
- The model incorrectly predicted that 46 customers would not churn but they actually churned

Random Forest has performed better than the logistic Regression model.

Hyperparameter tuning

```

In [46]: 1 # Define hyperparameter grid
2 from sklearn.model_selection import GridSearchCV
3 param_grid = {
4     'n_estimators': [100, 125, 150],
5     'max_depth': [5, 10, 15, 20, 25],
6     'min_samples_leaf': [1, 2],
7     'min_samples_split': [2, 5],
8     'criterion': ['gini', 'entropy'],
9 }
10
11 # Create Random Forest model
12 clf = RandomForestClassifier(random_state=0)
13
14 # Initialize GridSearchCV with early stopping and smaller training frac
15 grid_search = GridSearchCV(clf, param_grid, scoring='f1', cv=3, n_jobs=
16
17 # Fit the model on a smaller portion of training data (replace 0.8 with
18 grid_search.fit(X_train[:int(0.8 * len(X_train))], y_train[:int(0.8 * 1
19
20 # Get best parameters and model
21 best_params = grid_search.best_params_
22 best_model = grid_search.best_estimator_
23
24 # Evaluate on the testing set
25 y_pred = best_model.predict(X_test)
26
27 # Calculate evaluation metrics
28 accuracy = accuracy_score(y_test, y_pred)
29 f1 = f1_score(y_test, y_pred)
30
31 print(f"Accuracy: {accuracy}")
32 print(f"F1-score: {f1}")
33 print(f"Best Hyperparameters: {best_params}")

```

Accuracy: 0.9295352323838081

F1-score: 0.6967741935483871

Best Hyperparameters: {'criterion': 'gini', 'max_depth': 25, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 150}

- about 92.9% of the instances was correctly predicted
- 69.6% F1-score shows a relatively good precision and recall

3. DecisionTrees

```

In [47]: 1 from sklearn.tree import DecisionTreeClassifier

```

```

In [48]: 1 #Instantiate DecisionTreeClassifier
2 dt_clf = DecisionTreeClassifier(random_state=42)

```


```

In [49]: 1 #Fit on the training data
2 dt_clf.fit(X_train, y_train)
3
4 #predict on the test set
5 dt_y_pred = dt_clf.predict(X_test)

```



```
In [50]: 1 cm3 = confusion_matrix(y_test, dt_y_pred)
2 # Visualize the confusion matrix
3 sns.heatmap(cm3, annot=True, fmt="d", cmap='Blues',
4             xticklabels=['Predicted Negative', 'Predicted Positive'],
5             yticklabels=['Actual Negative', 'Actual Positive'])
6 plt.show()
```



- We can see that the model made correct predictions for 541 negative cases and 74 positive cases. However, it incorrectly predicted 25 positive cases as negative and 27 negative cases as positive.

```
In [51]: 1 # Calculate accuracy
2 accuracy = accuracy_score(y_test, dt_y_pred)
3 print(f"Accuracy: {accuracy}")
4
5 # Calculate F1-score
6 f1 = f1_score(y_test, dt_y_pred)
7 print(f"F1-score: {f1}")
8
9 # Calculate precision and recall for different thresholds
10 precision, recall, thresholds = precision_recall_curve(y_test, dt_y_pre
11
12 # Calculate ROC AUC
13 roc_auc = roc_auc_score(y_test, dt_y_pred)
14
15 print(f"ROC AUC: {roc_auc}")
```

Accuracy: 0.9220389805097451

F1-score: 0.74

ROC AUC: 0.844251828009656

4. XGBoost

```
In [52]: 1 from xgboost import XGBClassifier
```

```
In [53]: 1 #instantiate XGBClassifier
2 xg_clf = XGBClassifier(random_state=42)
3
4 #Fit on the training data
5 xg_clf.fit(X_train,y_train)
```

```
Out[53]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                        colsample_bylevel=None, colsample_bynode=None,
                        colsample_bytree=None, device=None, early_stopping_rounds=None,
                        enable_categorical=False, eval_metric=None, feature_types=None,
                        gamma=None, grow_policy=None, importance_type=None,
                        interaction_constraints=None, learning_rate=None, max_bin=None,
                        max_cat_threshold=None, max_cat_to_onehot=None,
                        max_delta_step=None, max_depth=None, max_leaves=None,
                        min_child_weight=None, missing=None, monotone_constraints=None,
                        multi_strategy=None, n_estimators=None, n_jobs=None,
                        num_parallel_tree=None, random_state=42, ...)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.


```
In [54]: 1 #predict on the test data
2 y_pred = xg_clf.predict(X_test)
```

In [55]:

```

1 cm2 = confusion_matrix(y_test, y_pred)
2 # Visualize the confusion matrix
3 sns.heatmap(cm2, annot=True, fmt="d", cmap='Greens',
4             xticklabels=['Predicted Negative', 'Predicted Positive'],
5             yticklabels=['Actual Negative', 'Actual Positive'])
6 plt.show()

```



- The matrix shows that the model correctly identified 562 negative cases (True Negatives) and 77 positive cases (True Positives). 4 negative cases were incorrectly predicted as positive (False Positives), and 24 positive cases were incorrectly predicted as negative (False Negatives).

In [56]:

```

1 # Calculate accuracy
2 accuracy = accuracy_score(y_test, y_pred)
3 print(f"Accuracy: {accuracy}")
4
5 # Calculate F1-score
6 f1 = f1_score(y_test, y_pred)
7 print(f"F1-score: {f1}")
8
9 # Calculate precision and recall for different thresholds
10 precision, recall, thresholds = precision_recall_curve(y_test, y_pred)
11
12 # Calculate ROC AUC
13 roc_auc = roc_auc_score(y_test, y_pred)
14
15 print(f"ROC AUC: {roc_auc}")

```

Accuracy: 0.9580209895052474

F1-score: 0.8461538461538463

ROC AUC: 0.8776545499072875

Show the AUC values of Logistic Regression, Random Forest models and XGBoost

```
In [57]: 1 # Make predictions on the test data
          2 y_pred_lr = logreg.predict(X_test)
          3
          4 # Calculate probability predictions for the positive class
          5 # Assuming positive class is at index 1
          6 lr_predictions_proba = logreg.predict_proba(X_test)[: , 1]

In [58]: 1 # Calculate probability predictions for the positive class
          2 clf.fit(X_train, y_train)
          3 rf_predictions_proba = clf.predict_proba(X_test)[: , 1]


In [59]: 1 # Calculate probability predictions for the positive class for theXGBoo
          2 xg_clf.fit(X_train, y_train)
          3 xg_predictions_proba = xg_clf.predict_proba(X_test)[: , 1]

In [60]: 1 # Calculate probability predictions for the positive class for decision
          2 dt_clf.fit(X_train, y_train)
          3 dt_predictions_proba = dt_clf.predict_proba(X_test)[: , 1]
```

```

In [61]: 1 # ROC curve calculation and plot
2 fpr_lr, tpr_lr, _ = roc_curve(y_test, lr_predictions_proba)
3 roc_auc_lr = auc(fpr_lr, tpr_lr)
4
5 fpr_rf, tpr_rf, _ = roc_curve(y_test, rf_predictions_proba)
6 roc_auc_rf = auc(fpr_rf, tpr_rf)
7
8 fpr_xg, tpr_xg, _ = roc_curve(y_test, xg_predictions_proba)
9 roc_auc_xg = auc(fpr_xg, tpr_xg)
10
11 fpr_dt, tpr_dt, _ = roc_curve(y_test, dt_predictions_proba)
12 roc_auc_dt = auc(fpr_dt, tpr_dt)
13
14 plt.figure(figsize=(8, 6))
15 plt.plot(fpr_lr, tpr_lr, label='Logistic Regression (AUC: {:.2f})'.format(roc_auc_lr),
16          linewidth=2)
17 plt.plot(fpr_rf, tpr_rf, label='Random Forest (AUC: {:.2f})'.format(roc_auc_rf),
18          linewidth=2)
19 plt.plot(fpr_xg, tpr_xg, label='XGBoost (AUC: {:.2f})'.format(roc_auc_xg),
20          linewidth=2)
21 plt.plot(fpr_dt, tpr_dt, label='Decision Tree (AUC: {:.2f})'.format(roc_auc_dt),
22          linewidth=2)
23
24 plt.plot([0, 1], [0, 1], linestyle='--', color='black', label='No Skill')
25 plt.xlabel('False Positive Rate (FPR)')
26 plt.ylabel('True Positive Rate (TPR)')
27 plt.title('ROC Curve Comparison')
28 plt.legend()
29 plt.grid(True)
30 plt.show()

```



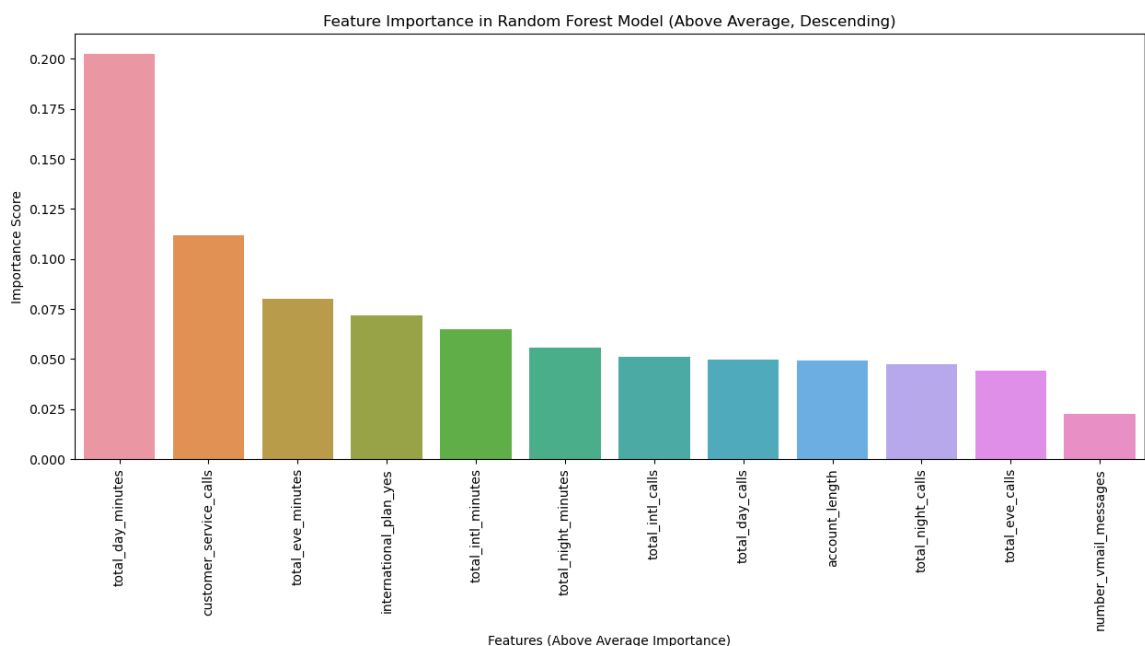
To check for important features

We use a simple bar graph to see the top features that can help determine churn rate

```
In [62]: 1 # first drop the 'state' columns
2 feature_names = df4.drop(columns=[col for col in df4.columns if 'state'
3 feature_names.columns
```

```
Out[62]: Index(['account_length', 'area_code', 'number_vmail_messages',
'total_day_minutes', 'total_day_calls', 'total_eve_minutes',
'total_eve_calls', 'total_night_minutes', 'total_night_calls',
'total_intl_minutes', 'total_intl_calls', 'customer_service_calls',
'churn', 'international_plan_yes', 'voice_mail_plan_yes'],
dtype='object')
```

```
In [63]: 1 # Set a larger figure size (adjust as needed)
2 feature_importances = clf.feature_importances_
3 feature_names = X_train.columns
4
5 # Calculate average importance score
6 avg_importance = feature_importances.mean()
7
8 # Filter features with above-average importance
9 important_features = feature_names[feature_importances > avg_importance]
10 important_importances = feature_importances[feature_importances > avg_i
11
12 # Sort features and importances in descending order
13 sorted_idx = important_importances.argsort()[::-1] # Reverse order for
14 sorted_features = important_features[sorted_idx]
15 sorted_importances = important_importances[sorted_idx]
16
17 # Create the bar plot
18 plt.figure(figsize=(15, 6))
19 sns.barplot(x=sorted_features, y=sorted_importances)
20 plt.xlabel('Features (Above Average Importance)')
21 plt.ylabel('Importance Score')
22 plt.title('Feature Importance in Random Forest Model (Above Average, De
23
24 # Rotate feature names for better readability
25 plt.xticks(rotation=90)
26
27 plt.show()
```



- From the graph, we can conclude the top features are:
 - Total day minutes
 - Customer service calls
 - Total eve minutes
 - Internal plan (Those that had subscribed)
 - Total minutes

Concusion

- The churn prediction analysis conducted for SyriaTel aimed to develop a classifier to identify customers likely to terminate their services. Through comprehensive data exploration, preparation, and modeling, several key findings emerged:
- Model Performance: Random Forest emerged as the most effective model for churn prediction, outperforming Logistic Regression, Decision Trees, and XGBoost. It exhibited superior accuracy and predictive power, making it the preferred choice for SyriaTel's churn prediction system.
- Key Predictive Features: Total day minutes, customer service calls, and subscription to the international plan were identified as crucial indicators of churn. These insights provide valuable guidance for SyriaTel in devising proactive retention strategies targeted at high-risk customers.#### Conclusion

Random Forest model appears to be the best model to predict the customers likely to churn.

Recommendation

Based on the plot, some recommendations would be:

- Enhance Call Quality: Invest in infrastructure and technology to improve call quality, ensuring a better customer experience.
- Customer Service Improvement: Focus on enhancing customer service by reducing response times, increasing efficiency in issue resolution, and offering personalized support.
- Tailored Plans for International Subscribers: Design attractive plans and offers specifically targeted at international subscribers to increase satisfaction and reduce churn.
- Proactive Retention Strategies: Implement proactive measures such as targeted promotions, loyalty rewards, and personalized communication to retain at-risk customers.
- Regular Analysis: Continuously monitor customer behavior and churn patterns, regularly updating models and strategies to adapt to changing market dynamics.